



✓ **Congratulations! You passed!**  
TO PASS 80% or higher

Keep Learning

GRADE  
100%

## Test prep

LATEST SUBMISSION GRADE

100%

1. Apache Spark pools in Azure Synapse Analytics benefit from which four of the following?

1 / 1 point

☒ Support for third party IDEs

✓ **Correct**

Azure Synapse Analytics provides an IDE for IntelliJ to create and submit applications to the Apache Spark pool.

☒ Autoscale

✓ **Correct**

Provide support for autoscale, so that pools can be scaled up or down as required, by adding or removing nodes.

☐ Real-time co-authoring (both authors see the changes in real-time)

☒ Pre-Loaded Anaconda libraries

✓ **Correct**

Over 200 Anaconda libraries are pre-installed on the Spark pool in Azure Synapse Analytics.

☒ REST APIs

✓ **Correct**

In order to monitor and submit jobs remotely, you can use Apache Livy as a Rest API for the Spark job server.

Apache Livy is a service that enables interaction with an Spark cluster over a REST interface to enable:

- submission of Spark jobs
- snippets of Spark code
- synchronous or asynchronous result retrieval
- and Spark Context management

2. Which one of the following resources should you create first when building an Apache Spark pool in Azure Synapse Analytics?

1 / 1 point

☐ A notebook

☒ A workspace

☐ A SQL Database

✓ **Correct**

An Azure Synapse Analytics workspace must be built before creating a new Apache Spark pool.

3. Which one of the following resources requires a unique Spark Pool name when creating an Apache Spark pool in Azure Synapse Analytics?

1 / 1 point

☒ The workspace

☐ The subscription

☐ The resource group

☐ Azure

✓ **Correct**

An Apache Spark pool name must adhere to the following specifications:

- Contain letters or numbers only,
- Contain 15 or less characters,
- Start with a letter,

- Contain no reserved words,
- and be unique in the workspace.

4. You need to create an embedded Spark capability that can reside alongside, and integrate with, other Azure data warehouse and integration services. Which one of the following solutions can provide these functionalities?

1 / 1 point

- ☒ Apache Spark for Azure Synapse
- ☐ Apache Spark
- ☐ Azure HDInsight
- ☐ Azure Databricks

✓ **Correct**

Apache Spark is an embedded Apache Spark capability within Azure Synapse Analytics residing on the same platform that contains data warehouses and data integration capabilities, as well as integrating with other Azure services.

5. You can use the magic command %%spark to execute which one of the following types of query against a spark context?

1 / 1 point

- ☒ Scala
- ☐ .Net for C#
- ☐ Python
- ☐ SparkSQL

✓ **Correct**

The magic command %%spark runs a Scala query against spark context.

6. You are using IDE-style IntelliSense to assist with your code completion. Which two of the following languages does the IntelliSense support for Syntax Code Completion?

1 / 1 point

- ☒ PySpark (Python)

✓ **Correct**

The IntelliSense supports Syntax Code Completion for PySpark (Python).

- ☐ .Net for Spark (C#)

- ☒ Spark (Scala)

✓ **Correct**

The IntelliSense supports Syntax Code Completion for Spark (Scala).

- ☐ SparkSQL

7. You run a notebook in Apache Spark pool for the first time. How many minutes will it take for Azure Synapse Analytics to create a new session?

1 / 1 point

- ☐ 30 - 60 Seconds
- ☐ 10 - 15 Minutes
- ☒ 3 - 5 Minutes

✓ **Correct**

The first time you run a notebook in a Spark pool, Synapse creates a new session. This can take approximately 3-5 minutes.

8. What is a DataFrame?

1 / 1 point

- ☐ A csv file
- ☐ A parquet file
- ☒ A data structure

✓ **Correct**

A DataFrame creates a data structure and it's one of the core data structures in Apache Spark.

9. You need to load data from different file types into an Apache Spark DataFrame. Which four of the following Azure services can you load the data from?

1 / 1 point

☒ Dedicated SQL pool

✓ **Correct**

You can load data into an Apache Spark DataFrame from different file types stored in an Azure Storage Account, or from data stored in a dedicated SQL pool.

☒ Azure Storage Account

✓ **Correct**

You can load data into an Apache Spark DataFrame from different file types stored in an Azure Storage Account, or from data stored in a dedicated SQL pool.

☒ Azure Data Lake Store Generation 2

✓ **Correct**

You can load data into an Apache Spark DataFrame from different file types stored in an Azure Storage Account, or from data stored in a dedicated SQL pool.

☐ Serverless SQL Pool

☒ Primary Storage Account

✓ **Correct**

You can load data into an Apache Spark DataFrame from different file types stored in an Azure Storage Account, or from data stored in a dedicated SQL pool.

10. You need to optimize the Apache Spark Jobs in Azure Synapse Analytics. However, first you need to consider the cluster configuration for the workload you're running on that cluster. Which one of the following resources can benefit from query optimization through Catalyst?

1 / 1 point

☒ DataFrames

☐ Resilient Distributed DataSets (RDDs)

☐ Notebooks

✓ **Correct**

DataFrames provide query optimization through Catalyst.